

# Advancing Weather Forecasting: A Comparative Analysis of Long Short-Term Memory and Support Vector Machines

## Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import seaborn as sns
from geopy.distance import distance
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf, month_plot, quarter_
from statsmodels.tsa.stattools import adfuller
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.preprocessing import MinMaxScaler
import torch
import torch.nn as nn
import time
from joblib import dump, load
```

## Loading the Dataset

```
In [2]: # Load the dataset
data = pd.read_csv('aggregated_hourly_data.csv')
data = data.dropna()
```

## Data Transformation

```
In [3]: # Calculate decimal degrees for Latitude and Longitude - even though there are geog
### but for this analysis i have not used geographical points
###though i have calculated some geographical variables

data['st_lat'] = data['st_lat'].apply(lambda x: int(str(x)[-2:])/3600 + int(str(x)[
data['st_long'] = data['st_long'].apply(lambda x: (int(str(x)[-2:])/3600 + int(str(
datahead = data.head(5)
datahead
```

Out[3]:

	Unnamed: 0	date	ind	rain	ind.1	temp	ind.2	wetb	dewpt	vappr	...	sun	vis
<b>370509</b>	370509	2007-12-31 02:00:00	2	0.0	0	9.6	0	8.8	8.0	10.7	...	0.0	20000.0
<b>370510</b>	370510	2007-12-31 03:00:00	0	0.0	0	9.7	0	8.8	7.9	10.6	...	0.0	20000.0
<b>370511</b>	370511	2007-12-31 04:00:00	0	0.0	0	9.9	0	8.8	7.6	10.4	...	0.0	20000.0
<b>370512</b>	370512	2007-12-31 05:00:00	0	0.0	0	10.4	0	9.3	8.2	10.8	...	0.0	20000.0
<b>370513</b>	370513	2007-12-31 06:00:00	0	0.0	0	10.6	0	9.7	8.7	11.3	...	0.0	20000.0

5 rows × 28 columns

```

In [4]: data.drop(columns=['Unnamed: 0'], inplace=True)
data.info()
datahead = data.head(100)
# rename columns
data.rename(columns = {'ind': 'i_rain', 'ind.1': 'i_temp', 'ind.2': 'i_wetb',
                        'ind.3': 'i_wdsp', 'ind.4': 'i_wddir'}, inplace =
print('Column names: ', data.columns)
# Are all datetime values unique?
print('Unique timestamps: ', data.date.nunique())
print('All timestamps: ', data.shape[0])
# Set timestamp as an index
data.set_index('date', inplace = True)
data.head()
# Convert index to datetime objects
data.index = pd.to_datetime(data.index)
# Check for missing hours in index
print('Unique Timestamps in our data: ', data.index.nunique())
print('Total range: ', (data.index.max() - data.index.min()) / pd.Timedelta('1 hour

```

```
<class 'pandas.core.frame.DataFrame'>
Index: 662345 entries, 370509 to 2962034
Data columns (total 27 columns):
```

#	Column	Non-Null Count	Dtype
0	date	662345 non-null	object
1	ind	662345 non-null	int64
2	rain	662345 non-null	float64
3	ind.1	662345 non-null	int64
4	temp	662345 non-null	float64
5	ind.2	662345 non-null	int64
6	wetb	662345 non-null	float64
7	dewpt	662345 non-null	float64
8	vappr	662345 non-null	float64
9	rhum	662345 non-null	float64
10	msl	662345 non-null	float64
11	ind.3	662345 non-null	float64
12	wdsp	662345 non-null	float64
13	ind.4	662345 non-null	float64
14	wddir	662345 non-null	float64
15	ww	662345 non-null	float64
16	w	662345 non-null	float64
17	sun	662345 non-null	float64
18	vis	662345 non-null	float64
19	clht	662345 non-null	float64
20	clamt	662345 non-null	float64
21	county	662345 non-null	object
22	st_id	662345 non-null	int64
23	st_name	662345 non-null	object
24	st_height	662345 non-null	int64
25	st_lat	662345 non-null	float64
26	st_long	662345 non-null	float64

```
dtypes: float64(19), int64(5), object(3)
```

```
memory usage: 141.5+ MB
```

```
Column names: Index(['date', 'i_rain', 'rain', 'i_temp', 'temp', 'i_wetb', 'wetb', 'dewpt', 'vappr', 'rhum', 'msl', 'i_wdsp', 'wdsp', 'i_wddir', 'wddir', 'ww', 'w', 'sun', 'vis', 'clht', 'clamt', 'county', 'st_id', 'st_name', 'st_height', 'st_lat', 'st_long'], dtype='object')
```

```
Unique timestamps: 123503
```

```
All timestamps: 662345
```

```
Unique Timestamps in our data: 123503
```

```
Total range: 123502.0
```

## Removing Nulls

```
In [5]: # count NAs
data.isnull().sum()
# Look at rows with NAs
data[data.isna().any(axis=1)]
# fill NAs by interpolating (default method = 'linear')
data.interpolate(inplace = True)
# make sure all NAs are filled
data.isnull().sum()
# drop indicators
data.drop(columns = ['i_rain', 'i_temp', 'i_wetb', 'i_wdsp', 'i_wddir'], inplace = True)
# count unique values in each column
data.nunique()
```

```
C:\Users\Syed Muhammad Mohsin\AppData\Local\Temp\ipykernel_25792\1707903422.py:6:  
FutureWarning: DataFrame.interpolate with object dtype is deprecated and will raise  
in a future version. Call obj.infer_objects(copy=False) before interpolating instead.  
data.interpolate(inplace = True)
```

```
Out[5]: rain      153  
temp      426  
wetb      342  
dewpt     327  
vappr     210  
rhum       80  
msl       978  
wdsp       53  
wddir      37  
ww         74  
w          71  
sun        11  
vis        87  
clht      157  
clamt      10  
county      6  
st_id       8  
st_name      8  
st_height    8  
st_lat       8  
st_long      8  
dtype: int64
```

## EXPLORATORY DATA ANALYSIS

```
In [6]: # Look at statistics  
Datadesc = data.describe().T  
#Datadesc.to_excel('datadesc.xlsx')  
Datadesc
```

Out[6]:


	count	mean	std	min	25%	50%	75%
rain	662345.0	0.124380	0.497221	0.000000	0.000000	0.000000	0.000000
temp	662345.0	9.922873	4.830016	-15.400000	6.600000	10.000000	13.400000
wetb	662345.0	8.596542	4.525487	-49.900000	5.500000	8.900000	11.900000
dewpt	662345.0	7.185166	4.513444	-16.400000	3.900000	7.400000	10.600000
vappr	662345.0	10.570086	3.175379	1.700000	8.100000	10.300000	12.800000
rhum	662345.0	84.142935	12.114841	20.000000	77.000000	87.000000	94.000000
msl	662345.0	1013.246276	12.529400	948.200000	1005.500000	1014.400000	1021.900000
wdsp	662345.0	9.918461	5.575393	0.000000	6.000000	9.000000	13.000000
wddir	662345.0	200.974477	84.909557	0.000000	140.000000	220.000000	260.000000
ww	662345.0	17.189089	23.643263	0.000000	2.000000	2.000000	25.000000
w	662345.0	33.520283	28.703269	0.000000	11.000000	11.000000	62.000000
sun	662345.0	0.167173	0.324983	0.000000	0.000000	0.000000	0.100000
vis	662345.0	26918.633144	15137.275498	5.000000	17000.000000	25000.000000	40000.000000
clht	662345.0	270.422165	401.669618	0.000000	20.000000	41.000000	250.000000
clamt	662345.0	5.746390	2.323820	0.000000	4.000000	7.000000	7.000000
st_id	662345.0	2477.993177	1668.160880	518.000000	532.000000	2375.000000	3904.000000
st_height	662345.0	86.559600	63.482789	9.000000	20.000000	71.000000	155.000000
st_lat	662345.0	53.042253	0.842819	51.847222	52.690278	53.305556	53.427222
st_long	662345.0	-7.964170	1.358236	-10.240833	-8.918056	-8.486111	-6.438889

### Correlation Matrix

In [7]:

```
# Select only numeric columns for correlation analysis - excluding geographical data
numeric_data = data.select_dtypes(include='number')
corr_matrix = numeric_data.corr()

# Look at correlations between weather variables recorded from Dublin Airport stations
sns.set_context('notebook', font_scale=1)
fig, ax = plt.subplots(figsize=(12, 12))
sns.heatmap(data=numeric_data.corr(), annot=True, fmt='.2f', cmap='coolwarm', ax=ax)
plt.show()
```



## Distribution of Variables - Transforming Data Weekly

```
In [8]: # making the weekly aggregated data-
## creating aggregated weekly data to analyse the temporal trends and patterns
### making it more effective than using raw data as an input

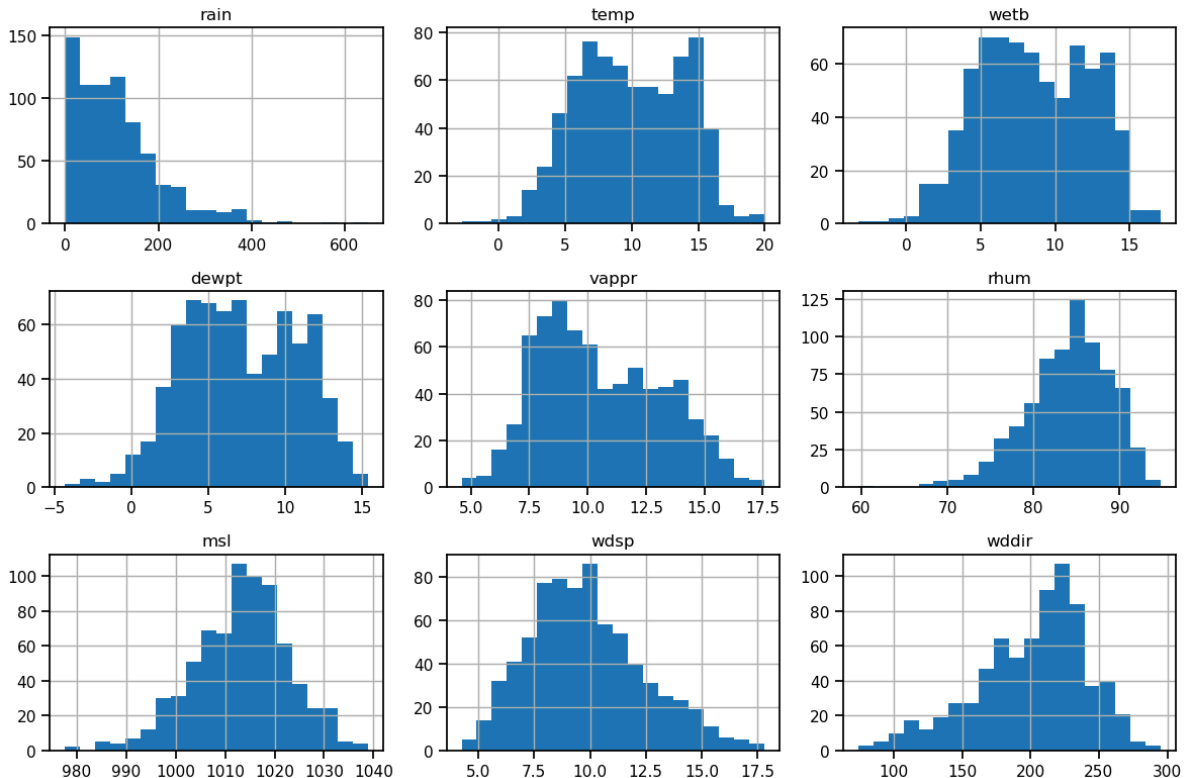
rainfall_weekly = data[['rain']].resample('W').sum()
all_weekly = data[['temp', 'wetb', 'dewpt', 'vaappr',
                  'rhum', 'msl', 'wdsp', 'wddir']].resample('W').mean()
weekly_da = pd.concat([rainfall_weekly, all_weekly], axis = 1)
weekly_da.isna().sum()

# visualize distribution of values for each variable with histogram
fig, axList = plt.subplots(3, 3, sharex=False, sharey=False)
axList = axList.flatten()
fig.set_size_inches(12, 8)
for i, ax in enumerate(axList[0:len(weekly_da.columns)]):
    weekly_da.hist(column = weekly_da.columns[i], bins = 20, ax=ax)
```

```
fig.tight_layout()
fig.show()
```

C:\Users\Syed Muhammad Mohsin\AppData\Local\Temp\ipykernel\_25792\2327037471.py:22: UserWarning: Matplotlib is currently using module://matplotlib\_inline.backend\_inline, which is a non-GUI backend, so cannot show the figure.

```
fig.show()
```



## MODELLING

### LSTM VS SVM

### LONG-SHORT TERM MEMORY - RAINFALL FORECASTING

```
In [9]: # Extract rainfall data - from the weekly data aggregated data created earlier
rainfall_data = weekly_da['rain'].values.astype(float)
# Normalize data - min max scaler
scaler = MinMaxScaler(feature_range=(-1, 1))
rainfall_data_normalized = scaler.fit_transform(rainfall_data.reshape(-1, 1))
# Convert data to PyTorch tensors
rainfall_tensor = torch.FloatTensor(rainfall_data_normalized).view(-1)
def create_sequences(data, seq_length):
    xs = []
    ys = []
    for i in range(len(data) - seq_length):
        x = data[i:i+seq_length]
        y = data[i+seq_length:i+seq_length+1]
        xs.append(x)
        ys.append(y)
    return torch.stack(xs), torch.stack(ys)

# Define sequence length
seq_length = 5
# Create sequences
xs, ys = create_sequences(rainfall_tensor, seq_length)
```

```

# Split data into training and testing sets
train_size = int(len(xs) * 0.8)
test_size = len(xs) - train_size
train_X, test_X = xs[:train_size], xs[train_size:]
train_y, test_y = ys[:train_size], ys[train_size:]

# Define LSTM model
class LSTM(nn.Module):
    def __init__(self, input_size=1, hidden_layer_size=100, output_size=1):
        super().__init__()
        self.hidden_layer_size = hidden_layer_size
        self.lstm = nn.LSTM(input_size, hidden_layer_size)
        self.linear = nn.Linear(hidden_layer_size, output_size)
        self.hidden_cell = (torch.zeros(1,1,self.hidden_layer_size),
                             torch.zeros(1,1,self.hidden_layer_size))

    def forward(self, input_seq):
        lstm_out, self.hidden_cell = self.lstm(input_seq.view(len(input_seq) ,1, -1))
        predictions = self.linear(lstm_out.view(len(input_seq), -1))
        return predictions[-1]

# Hyperparameters for tuning - could have used more hidden layers but computational
learning_rates = [0.001, 0.01]
hidden_layer_sizes = [10, 50]

# Record validation losses and MSEs and Record training times
validation_losses = []
mse_values = []
training_times_rain_lstm = []

# Train models with different hyperparameters
for lr in learning_rates:
    for hidden_size in hidden_layer_sizes:
        start_time = time.time() # started time for collecting training times
        # Instantiate the model, define loss function and optimizer
        model = LSTM(hidden_layer_size=hidden_size)
        loss_function = nn.MSELoss()
        optimizer = torch.optim.Adam(model.parameters(), lr=lr)

        # Train the model with epochs set to 50
        epochs = 50
        train_losses = []
        val_losses = []

        for i in range(epochs):
            model.train()
            epoch_train_loss = 0
            # Record start time
            start_time = time.time()
            for seq, labels in zip(train_X, train_y):
                optimizer.zero_grad()
                model.hidden_cell = (torch.zeros(1, 1, model.hidden_layer_size),
                                     torch.zeros(1, 1, model.hidden_layer_size))

                y_pred = model(seq)

                single_loss = loss_function(y_pred, labels)
                epoch_train_loss += single_loss.item()
                single_loss.backward()
                optimizer.step()

            # Calculate epoch training time and appending epoch training time to the
            epoch_time = time.time() - start_time

```



```

training_times_rain_lstm.append(epoch_time)
train_losses.append(epoch_train_loss / len(train_X))

model.eval()
epoch_val_loss = 0
with torch.no_grad():
    for seq, labels in zip(test_X, test_y):
        model.hidden_cell = (torch.zeros(1, 1, model.hidden_layer_size),
                               torch.zeros(1, 1, model.hidden_layer_size))

        y_pred = model(seq)
        val_loss = loss_function(y_pred, labels)
        epoch_val_loss += val_loss.item()

    # Record validation Loss
    val_losses.append(epoch_val_loss / len(test_X))

if i % 10 == 1:
    print(f'lr: {lr}, hidden_size: {hidden_size}, epoch: {i:3} train_loss: {train_losses[-1]:.4f}, validation_loss: {val_losses[-1]:.4f}')

validation_losses.append(val_losses)

# Make predictions using the trained model lstm-rainfall
model.eval()
test_predictions = []
for seq in test_X:
    with torch.no_grad():
        model.hidden = (torch.zeros(1, 1, model.hidden_layer_size),
                          torch.zeros(1, 1, model.hidden_layer_size))
        test_predictions.append(model(seq).item())

# Inverse transform the predictions
test_predictions = scaler.inverse_transform(np.array(test_predictions).reshape(-1, 1))
# Calculate MSE for rainfall
mse_rainfall = ((scaler.inverse_transform(test_y.reshape(-1, 1)) - test_predictions)**2).mean()
mse_values.append({'lr': lr,
                   'hidden_size': hidden_size,
                   'mse_rainfall': mse_rainfall})

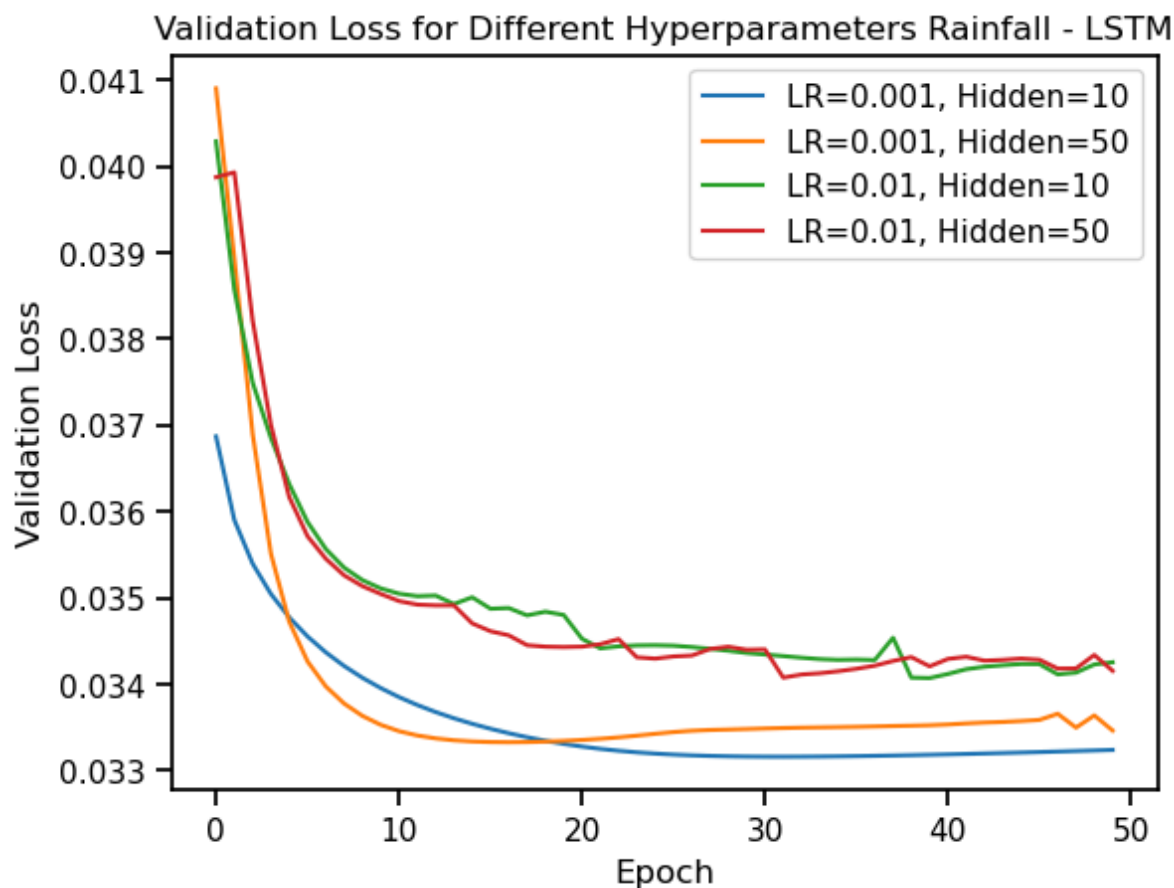
# Plot validation Loss graphs for each set of hyperparameters
for idx, (lr, hidden_size) in enumerate([(lr, hidden_size) for lr in learning_rates for hidden_size in hidden_sizes]):
    plt.plot(validation_losses[idx], label=f'LR={lr}, Hidden={hidden_size}')

plt.xlabel('Epoch')
plt.ylabel('Validation Loss')
plt.title('Validation Loss for Different Hyperparameters Rainfall - LSTM')
plt.legend()
plt.show()

# Convert MSE values to DataFrame
mse_df_rain_lstm = pd.DataFrame(mse_values)
# Display the MSE values table
print("MSE Values Rainfall:")
print(mse_df_rain_lstm)

```

lr: 0.001, hidden\_size: 10, epoch: 1 train\_loss\_rainfall: 0.07602387 val\_loss\_rainfall: 0.03590120 epoch\_time: 1.25 seconds  
lr: 0.001, hidden\_size: 10, epoch: 11 train\_loss\_rainfall: 0.07180347 val\_loss\_rainfall: 0.03375670 epoch\_time: 1.15 seconds  
lr: 0.001, hidden\_size: 10, epoch: 21 train\_loss\_rainfall: 0.07109370 val\_loss\_rainfall: 0.03324873 epoch\_time: 1.23 seconds  
lr: 0.001, hidden\_size: 10, epoch: 31 train\_loss\_rainfall: 0.07083794 val\_loss\_rainfall: 0.03315763 epoch\_time: 1.20 seconds  
lr: 0.001, hidden\_size: 10, epoch: 41 train\_loss\_rainfall: 0.07064841 val\_loss\_rainfall: 0.03318968 epoch\_time: 1.27 seconds  
lr: 0.001, hidden\_size: 50, epoch: 1 train\_loss\_rainfall: 0.07680452 val\_loss\_rainfall: 0.03893052 epoch\_time: 1.51 seconds  
lr: 0.001, hidden\_size: 50, epoch: 11 train\_loss\_rainfall: 0.07139355 val\_loss\_rainfall: 0.03340363 epoch\_time: 1.43 seconds  
lr: 0.001, hidden\_size: 50, epoch: 21 train\_loss\_rainfall: 0.07077269 val\_loss\_rainfall: 0.03336529 epoch\_time: 1.46 seconds  
lr: 0.001, hidden\_size: 50, epoch: 31 train\_loss\_rainfall: 0.06968960 val\_loss\_rainfall: 0.03349022 epoch\_time: 1.56 seconds  
lr: 0.001, hidden\_size: 50, epoch: 41 train\_loss\_rainfall: 0.06886658 val\_loss\_rainfall: 0.03354351 epoch\_time: 1.47 seconds  
lr: 0.01, hidden\_size: 10, epoch: 1 train\_loss\_rainfall: 0.07623863 val\_loss\_rainfall: 0.03857206 epoch\_time: 1.15 seconds  
lr: 0.01, hidden\_size: 10, epoch: 11 train\_loss\_rainfall: 0.07295657 val\_loss\_rainfall: 0.03501853 epoch\_time: 1.15 seconds  
lr: 0.01, hidden\_size: 10, epoch: 21 train\_loss\_rainfall: 0.07237298 val\_loss\_rainfall: 0.03441368 epoch\_time: 1.14 seconds  
lr: 0.01, hidden\_size: 10, epoch: 31 train\_loss\_rainfall: 0.07149480 val\_loss\_rainfall: 0.03432428 epoch\_time: 1.13 seconds  
lr: 0.01, hidden\_size: 10, epoch: 41 train\_loss\_rainfall: 0.07100196 val\_loss\_rainfall: 0.03417190 epoch\_time: 1.11 seconds  
lr: 0.01, hidden\_size: 50, epoch: 1 train\_loss\_rainfall: 0.07754331 val\_loss\_rainfall: 0.03993025 epoch\_time: 1.49 seconds  
lr: 0.01, hidden\_size: 50, epoch: 11 train\_loss\_rainfall: 0.07268040 val\_loss\_rainfall: 0.03492030 epoch\_time: 1.41 seconds  
lr: 0.01, hidden\_size: 50, epoch: 21 train\_loss\_rainfall: 0.07207873 val\_loss\_rainfall: 0.03446159 epoch\_time: 1.54 seconds  
lr: 0.01, hidden\_size: 50, epoch: 31 train\_loss\_rainfall: 0.07323821 val\_loss\_rainfall: 0.03407450 epoch\_time: 1.48 seconds  
lr: 0.01, hidden\_size: 50, epoch: 41 train\_loss\_rainfall: 0.06852092 val\_loss\_rainfall: 0.03431610 epoch\_time: 1.43 seconds



MSE Values Rainfall:

	lr	hidden_size	mse_rainfall
0	0.001	10	3550.368425
1	0.001	50	3533.907725
2	0.010	10	4174.129595
3	0.010	50	3958.486429

## Creating a df to save model losses ( train and validation) for comparison

```
In [10]: # Create a DataFrame for hyperparameters, validation losses, and training times for
lstm_data = []
for lr in learning_rates:
    for hidden_size in hidden_layer_sizes:
        idx = learning_rates.index(lr) * len(hidden_layer_sizes) + hidden_layer_size
        lstm_data.append({
            'Learning Rate': lr,
            'Hidden Size': hidden_size,
            'Validation Loss': validation_losses[idx][-1], # Get the last validation loss
            'Training Time': training_times_rain_lstm[idx] # Get the training time
        })

# Convert the LSTM data to a DataFrame
df_hyperparameters_lstm_rain = pd.DataFrame(lstm_data)

# Display the DataFrame
print(df_hyperparameters_lstm_rain)
```

	Learning Rate	Hidden Size	Validation Loss	Training Time
0	0.001	10	0.033237	1.090688
1	0.001	50	0.033461	1.254242
2	0.010	10	0.034252	1.224217
3	0.010	50	0.034153	1.120978

```
In [11]: #This line of code is generated to draw the hyperparameter comparison (using a df to
#df_hyperparameters_lstm_rain.to_excel('df_hyperparameters_lstm_rain.xlsx')
```

## Best LSTM Model Trained with best hyperparameters - Rainfall

```
In [12]: # Find the index of the minimum validation Loss
best_index = np.unravel_index(np.argmin(validation_losses), np.array(validation_losses).shape)
best_lr_index = best_index[0] // len(hidden_layer_sizes)
best_hidden_size_index = best_index[0] % len(hidden_layer_sizes)
best_lr = learning_rates[best_lr_index]
best_hidden_size = hidden_layer_sizes[best_hidden_size_index]

# Instantiate the best model with the best hyperparameters
best_model_lstm_rainfall = LSTM(hidden_layer_size=best_hidden_size)
loss_function = nn.MSELoss()
optimizer = torch.optim.Adam(best_model_lstm_rainfall.parameters(), lr=best_lr)

# Train the best model
## Train the best model lstm rainfall over best set of hyperparameters deduced early
epochs = 100
best_train_losses = []
best_val_losses = []

for i in range(epochs):
    best_model_lstm_rainfall.train()
    epoch_train_loss = 0
    for seq, labels in zip(train_X, train_y):
        optimizer.zero_grad()
        best_model_lstm_rainfall.hidden_cell = (torch.zeros(1, 1, best_model_lstm_rainfall.hidden_layer_size),
                                                torch.zeros(1, 1, best_model_lstm_rainfall.hidden_layer_size))

        y_pred = best_model_lstm_rainfall(seq)

        single_loss = loss_function(y_pred, labels)
        epoch_train_loss += single_loss.item()
        single_loss.backward()
        optimizer.step()

    best_train_losses.append(epoch_train_loss / len(train_X))

    best_model_lstm_rainfall.eval()
    epoch_val_loss = 0
    with torch.no_grad():
        for seq, labels in zip(test_X, test_y):
            best_model_lstm_rainfall.hidden_cell = (torch.zeros(1, 1, best_model_lstm_rainfall.hidden_layer_size),
                                                    torch.zeros(1, 1, best_model_lstm_rainfall.hidden_layer_size))

            y_pred = best_model_lstm_rainfall(seq)
            val_loss = loss_function(y_pred, labels)
            epoch_val_loss += val_loss.item()

    # Record validation Loss
    best_val_losses.append(epoch_val_loss / len(test_X))

    if i % 25 == 1:
        print(f'Best model: lr={best_lr}, hidden_size={best_hidden_size}, epoch: {i}')

# Make predictions using the best model
best_model_lstm_rainfall.eval()
best_test_predictions_rain_lstm = []
for seq in test_X:
    with torch.no_grad():
```

```

best_model_lstm_rainfall.hidden = (torch.zeros(1, 1, best_model_lstm_rainfa
    torch.zeros(1, 1, best_model_lstm_rainfall.hidden_layer_siz
best_test_predictions_rain_lstm.append(best_model_lstm_rainfall(seq).item())

# Inverse transform the predictions
best_test_predictions_rain_lstm = scaler.inverse_transform(np.array(best_test_predi

# Plot actual vs predicted for the best model
plt.plot(best_test_predictions_rain_lstm, label='Predicted')
plt.plot(scaler.inverse_transform(test_y.reshape(-1, 1)), label='Actual')
plt.xlabel('Time')
plt.ylabel('Rainfall')
plt.title('Actual vs Predicted Rainfall for Best Model LSTM')
plt.legend()
plt.show()

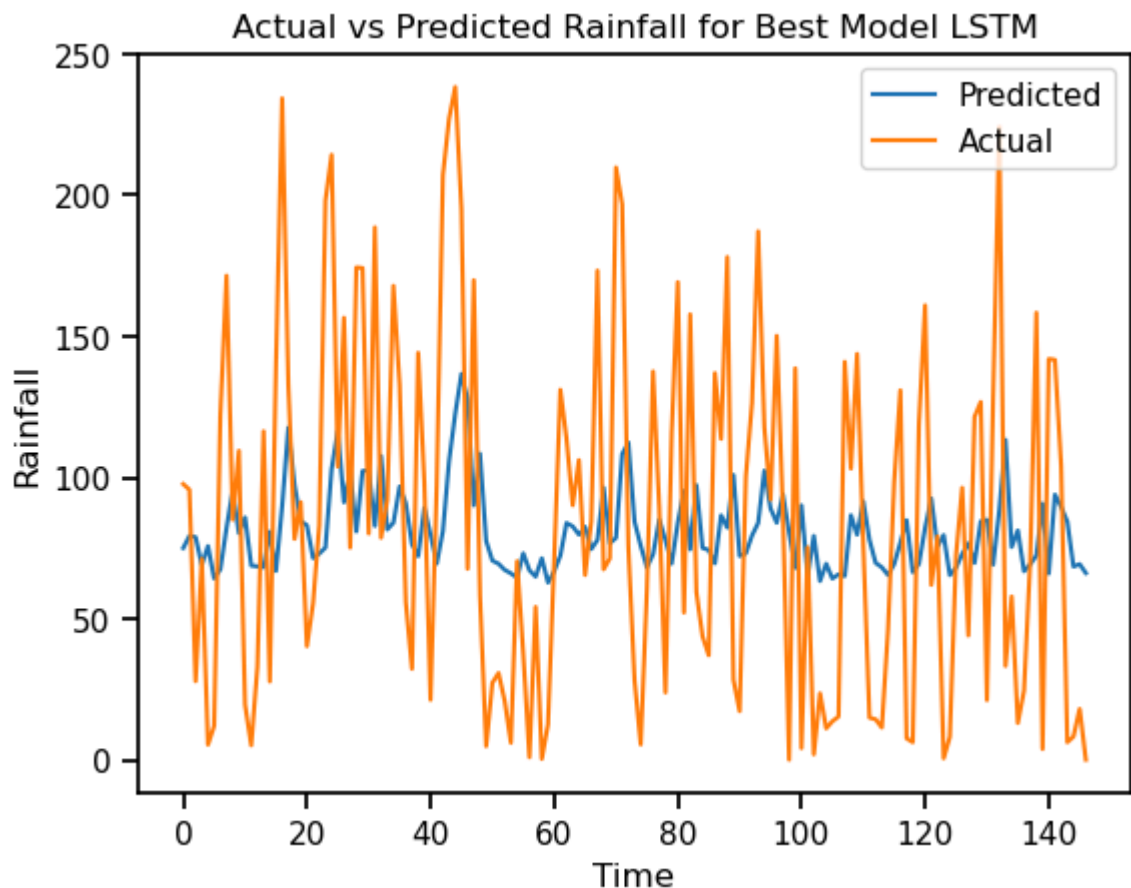
```

Best model: lr=0.001, hidden\_size=10, epoch: 1 train\_loss\_rainfall: 0.07630340 v  
al\_loss\_rainfall: 0.03624599

Best model: lr=0.001, hidden\_size=10, epoch: 26 train\_loss\_rainfall: 0.07075860 v  
al\_loss\_rainfall: 0.03347538

Best model: lr=0.001, hidden\_size=10, epoch: 51 train\_loss\_rainfall: 0.06944639 v  
al\_loss\_rainfall: 0.03419261

Best model: lr=0.001, hidden\_size=10, epoch: 76 train\_loss\_rainfall: 0.06772627 v  
al\_loss\_rainfall: 0.03409336



## Saving the test file

```

In [13]: # Save the test data as a NumPy file
np.save('test_X.npy', test_X)

```

## Saving the Model-LSTM\_Rainfall

```
In [14]: #### SAVING THE MODEL- LSTM
# Save the best model as a joblib file
dump(best_model_lstm_rainfall, 'best_model_lstm_rainfall.joblib')

# Load the best model from the joblib file ( if needed)
#loaded_model = load('best_model_lstm_rainfall.joblib')

# Load the test data from the NumPy file
#loaded_test_X = np.load('test_X.npy')

## to use this test text convert to the test to pytorch tensor to be utilized
```

Out[14]: ['best\_model\_lstm\_rainfall.joblib']

## LONG-SHORT TERM MEMORY - TEMPERATURE FORECASTING

```
In [15]: # Extract temperature data
temperature_data = weekly_da['temp'].values.astype(float)

# Normalize data
scaler = MinMaxScaler(feature_range=(-1, 1))
temperature_data_normalized = scaler.fit_transform(temperature_data.reshape(-1, 1))

# Convert data to PyTorch tensors
temperature_tensor = torch.FloatTensor(temperature_data_normalized).view(-1)

def create_sequences(data, seq_length):
    xs = []
    ys = []
    for i in range(len(data) - seq_length):
        x = data[i:i+seq_length]
        y = data[i+seq_length:i+seq_length+1]
        xs.append(x)
        ys.append(y)
    return torch.stack(xs), torch.stack(ys)

# Define sequence length
seq_length = 5

# Create sequences
xs, ys = create_sequences(temperature_tensor, seq_length)

# Split data into training and testing sets
train_size = int(len(xs) * 0.8)
test_size = len(xs) - train_size
train_X, test_X = xs[:train_size], xs[train_size:]
train_y, test_y = ys[:train_size], ys[train_size:]

# Define LSTM model
class LSTM(nn.Module):
    def __init__(self, input_size=1, hidden_layer_size=100, output_size=1):
        super().__init__()
        self.hidden_layer_size = hidden_layer_size
        self.lstm = nn.LSTM(input_size, hidden_layer_size)
        self.linear = nn.Linear(hidden_layer_size, output_size)
        self.hidden_cell = (torch.zeros(1,1,self.hidden_layer_size),
                             torch.zeros(1,1,self.hidden_layer_size))
```

```

def forward(self, input_seq):
    lstm_out, self.hidden_cell = self.lstm(input_seq.view(len(input_seq) ,1, -1)
    predictions = self.linear(lstm_out.view(len(input_seq), -1))
    return predictions[-1]

# Hyperparameters for tuning
learning_rates = [0.001, 0.01]
hidden_layer_sizes = [10, 50]

# Record validation losses and MSEs
validation_losses = []
mse_values = []
# Record training times
training_times_temp_lstm = []

# Train models with different hyperparameters
for lr in learning_rates:
    for hidden_size in hidden_layer_sizes:
        # Instantiate the model, define loss function and optimizer
        model = LSTM(hidden_layer_size=hidden_size)
        loss_function = nn.MSELoss()
        optimizer = torch.optim.Adam(model.parameters(), lr=lr)

        # Train the model
        epochs = 50
        train_losses = []
        val_losses = []

        for i in range(epochs):
            model.train()
            epoch_train_loss = 0
            # Record start time
            start_time = time.time()
            for seq, labels in zip(train_X, train_y):
                optimizer.zero_grad()
                model.hidden_cell = (torch.zeros(1, 1, model.hidden_layer_size),
                                     torch.zeros(1, 1, model.hidden_layer_size))

                y_pred = model(seq)

                single_loss = loss_function(y_pred, labels)
                epoch_train_loss += single_loss.item()
                single_loss.backward()
                optimizer.step()

            # Calculate epoch training time
            epoch_time = time.time() - start_time

            # Append epoch training time to the list
            training_times_temp_lstm.append(epoch_time)
            train_losses.append(epoch_train_loss / len(train_X))

            model.eval()
            epoch_val_loss = 0
            with torch.no_grad():
                for seq, labels in zip(test_X, test_y):
                    model.hidden_cell = (torch.zeros(1, 1, model.hidden_layer_size),
                                         torch.zeros(1, 1, model.hidden_layer_size))

                    y_pred = model(seq)
                    val_loss = loss_function(y_pred, labels)
                    epoch_val_loss += val_loss.item()

            # Record validation loss

```

```

        val_losses.append(epoch_val_loss / len(test_X))

    if i % 10 == 1:
        print(f'lr: {lr}, hidden_size: {hidden_size}, epoch: {i:3} train_loss: {train_loss}')

    validation_losses.append(val_losses)

    # Make predictions using the trained model
    model.eval()
    test_predictions = []
    for seq in test_X:
        with torch.no_grad():
            model.hidden = (torch.zeros(1, 1, model.hidden_layer_size),
                           torch.zeros(1, 1, model.hidden_layer_size))
            test_predictions.append(model(seq).item())

    # Inverse transform the predictions
    test_predictions = scaler.inverse_transform(np.array(test_predictions).reshape(-1, 1))

    # Calculate MSE for temperature
    mse_temperature = ((scaler.inverse_transform(test_y.reshape(-1, 1)) - test_predictions)**2).mean()

    mse_values.append({
        'lr': lr,
        'hidden_size': hidden_size,
        'mse_temperature': mse_temperature
    })

# Plot validation loss graphs for each set of hyperparameters
for idx, (lr, hidden_size) in enumerate([(lr, hidden_size) for lr in learning_rates for hidden_size in hidden_sizes]):
    plt.plot(validation_losses[idx], label=f'LR={lr}, Hidden={hidden_size}')

plt.xlabel('Epoch')
plt.ylabel('Validation Loss')
plt.title('Validation Loss for Different Hyperparameters Temperature - LSTM')
plt.legend()
plt.show()

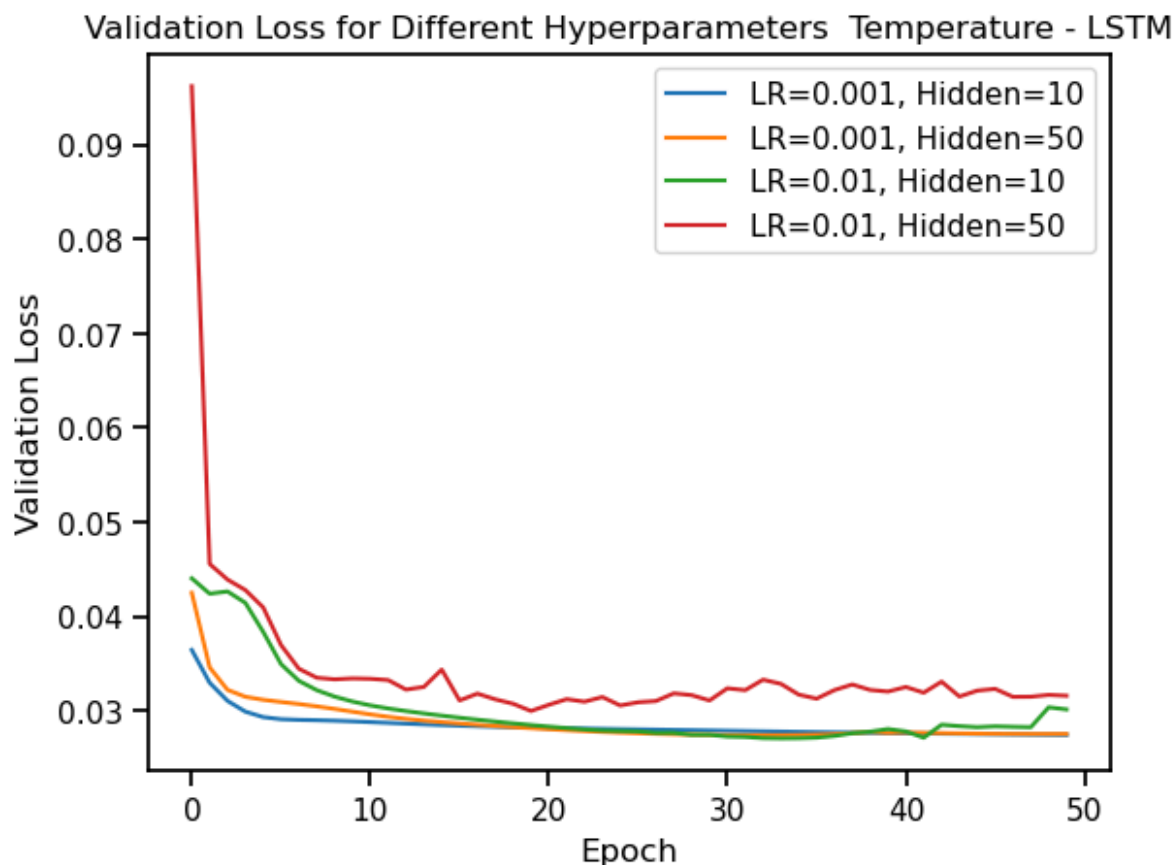
# Convert MSE values to DataFrame
mse_df_temp_lstm = pd.DataFrame(mse_values)

# Display the MSE values table
print("MSE Values:")
print(mse_df_temp_lstm)

```



lr: 0.001, hidden\_size: 10, epoch: 1 train\_loss\_temperature: 0.03826221 val\_loss\_temperature: 0.03296080 epoch\_time: 1.16 seconds  
lr: 0.001, hidden\_size: 10, epoch: 11 train\_loss\_temperature: 0.02787560 val\_loss\_temperature: 0.02870887 epoch\_time: 1.15 seconds  
lr: 0.001, hidden\_size: 10, epoch: 21 train\_loss\_temperature: 0.02729108 val\_loss\_temperature: 0.02813397 epoch\_time: 1.13 seconds  
lr: 0.001, hidden\_size: 10, epoch: 31 train\_loss\_temperature: 0.02702738 val\_loss\_temperature: 0.02783833 epoch\_time: 1.14 seconds  
lr: 0.001, hidden\_size: 10, epoch: 41 train\_loss\_temperature: 0.02683518 val\_loss\_temperature: 0.02758594 epoch\_time: 1.16 seconds  
lr: 0.001, hidden\_size: 50, epoch: 1 train\_loss\_temperature: 0.03464158 val\_loss\_temperature: 0.03459286 epoch\_time: 1.41 seconds  
lr: 0.001, hidden\_size: 50, epoch: 11 train\_loss\_temperature: 0.02823700 val\_loss\_temperature: 0.02934927 epoch\_time: 1.39 seconds  
lr: 0.001, hidden\_size: 50, epoch: 21 train\_loss\_temperature: 0.02730909 val\_loss\_temperature: 0.02792965 epoch\_time: 1.43 seconds  
lr: 0.001, hidden\_size: 50, epoch: 31 train\_loss\_temperature: 0.02672229 val\_loss\_temperature: 0.02739959 epoch\_time: 1.40 seconds  
lr: 0.001, hidden\_size: 50, epoch: 41 train\_loss\_temperature: 0.02593279 val\_loss\_temperature: 0.02765058 epoch\_time: 1.39 seconds  
lr: 0.01, hidden\_size: 10, epoch: 1 train\_loss\_temperature: 0.03687944 val\_loss\_temperature: 0.04240744 epoch\_time: 1.14 seconds  
lr: 0.01, hidden\_size: 10, epoch: 11 train\_loss\_temperature: 0.03220438 val\_loss\_temperature: 0.03023518 epoch\_time: 1.29 seconds  
lr: 0.01, hidden\_size: 10, epoch: 21 train\_loss\_temperature: 0.03141422 val\_loss\_temperature: 0.02815167 epoch\_time: 1.16 seconds  
lr: 0.01, hidden\_size: 10, epoch: 31 train\_loss\_temperature: 0.03068597 val\_loss\_temperature: 0.02721832 epoch\_time: 1.14 seconds  
lr: 0.01, hidden\_size: 10, epoch: 41 train\_loss\_temperature: 0.03100885 val\_loss\_temperature: 0.02713575 epoch\_time: 1.16 seconds  
lr: 0.01, hidden\_size: 50, epoch: 1 train\_loss\_temperature: 0.04605461 val\_loss\_temperature: 0.04556071 epoch\_time: 1.36 seconds  
lr: 0.01, hidden\_size: 50, epoch: 11 train\_loss\_temperature: 0.03386287 val\_loss\_temperature: 0.03324554 epoch\_time: 1.36 seconds  
lr: 0.01, hidden\_size: 50, epoch: 21 train\_loss\_temperature: 0.03110022 val\_loss\_temperature: 0.03121962 epoch\_time: 1.37 seconds  
lr: 0.01, hidden\_size: 50, epoch: 31 train\_loss\_temperature: 0.02798713 val\_loss\_temperature: 0.03217605 epoch\_time: 1.41 seconds  
lr: 0.01, hidden\_size: 50, epoch: 41 train\_loss\_temperature: 0.02634635 val\_loss\_temperature: 0.03190569 epoch\_time: 1.42 seconds



MSE Values:

	lr	hidden_size	mse_temperature
0	0.001	10	3.741914
1	0.001	50	3.683452
2	0.010	10	4.937453
3	0.010	50	5.029695

## Creating a df to save model losses ( train and validation) for comparison

```
In [16]: # Create a DataFrame for hyperparameters, validation losses, and training times for
lstm_data_temp = []
for lr in learning_rates:
    for hidden_size in hidden_layer_sizes:
        idx = learning_rates.index(lr) * len(hidden_layer_sizes) + hidden_layer_sizes.index(hidden_size)
        lstm_data_temp.append({
            'Learning Rate': lr,
            'Hidden Size': hidden_size,
            'Validation Loss': validation_losses[idx][-1], # Get the last validation loss
            'Training Time': training_times_rain_lstm[idx] # Get the training time
        })

# Convert the LSTM data to a DataFrame
df_hyperparameters_lstm_temp = pd.DataFrame(lstm_data_temp)

# Display the DataFrame
print(df_hyperparameters_lstm_temp)
```

	Learning Rate	Hidden Size	Validation Loss	Training Time
0	0.001	10	0.027431	1.090688
1	0.001	50	0.027529	1.254242
2	0.010	10	0.030135	1.224217
3	0.010	50	0.031584	1.120978

```
In [17]: #This line of code is generated to draw the hyperparameter comparison (using a df to
#df_hyperparameters_lstm_temp.to_excel('df_hyperparameters_lstm_temp.xlsx')
```

## Best LSTM Model Trained with best hyperparameters - Temperature

```
In [18]: # Find the index of the minimum validation loss
best_index = np.unravel_index(np.argmin(validation_losses), np.array(validation_losses).shape)
best_lr_index = best_index[0] // len(hidden_layer_sizes)
best_hidden_size_index = best_index[0] % len(hidden_layer_sizes)
best_lr = learning_rates[best_lr_index]
best_hidden_size = hidden_layer_sizes[best_hidden_size_index]

# Instantiate the best model with the best hyperparameters
best_model_lstm_temp = LSTM(hidden_layer_size=best_hidden_size)
loss_function = nn.MSELoss()
optimizer = torch.optim.Adam(best_model_lstm_temp.parameters(), lr=best_lr)

# Train the best model lstm temperature over best set of hyperparameters deduced epochs
epochs = 100
best_train_losses = [] # calculating training losses and validation losses
best_val_losses = []

for i in range(epochs):
    best_model_lstm_temp.train()
    epoch_train_loss = 0
    for seq, labels in zip(train_X, train_y):
        optimizer.zero_grad()
        best_model_lstm_temp.hidden_cell = (torch.zeros(1, 1, best_model_lstm_temp.hidden_layer_size),
                                             torch.zeros(1, 1, best_model_lstm_temp.hidden_layer_size))

        y_pred = best_model_lstm_temp(seq)

        single_loss = loss_function(y_pred, labels)
        epoch_train_loss += single_loss.item()
        single_loss.backward()
        optimizer.step()

    best_train_losses.append(epoch_train_loss / len(train_X))

    best_model_lstm_temp.eval()
    epoch_val_loss = 0
    with torch.no_grad():
        for seq, labels in zip(test_X, test_y):
            best_model_lstm_temp.hidden_cell = (torch.zeros(1, 1, best_model_lstm_temp.hidden_layer_size),
                                                torch.zeros(1, 1, best_model_lstm_temp.hidden_layer_size))

            y_pred = best_model_lstm_temp(seq)
            val_loss = loss_function(y_pred, labels)
            epoch_val_loss += val_loss.item()

    # Record validation loss
    best_val_losses.append(epoch_val_loss / len(test_X))

    if i % 25 == 1:
        print(f'Best model: lr={best_lr}, hidden_size={best_hidden_size}, epoch: {i}')

# Make predictions using the best model
best_model_lstm_temp.eval()
best_test_predictions_lstm_temp = []
for seq in test_X:
    with torch.no_grad():
```

```

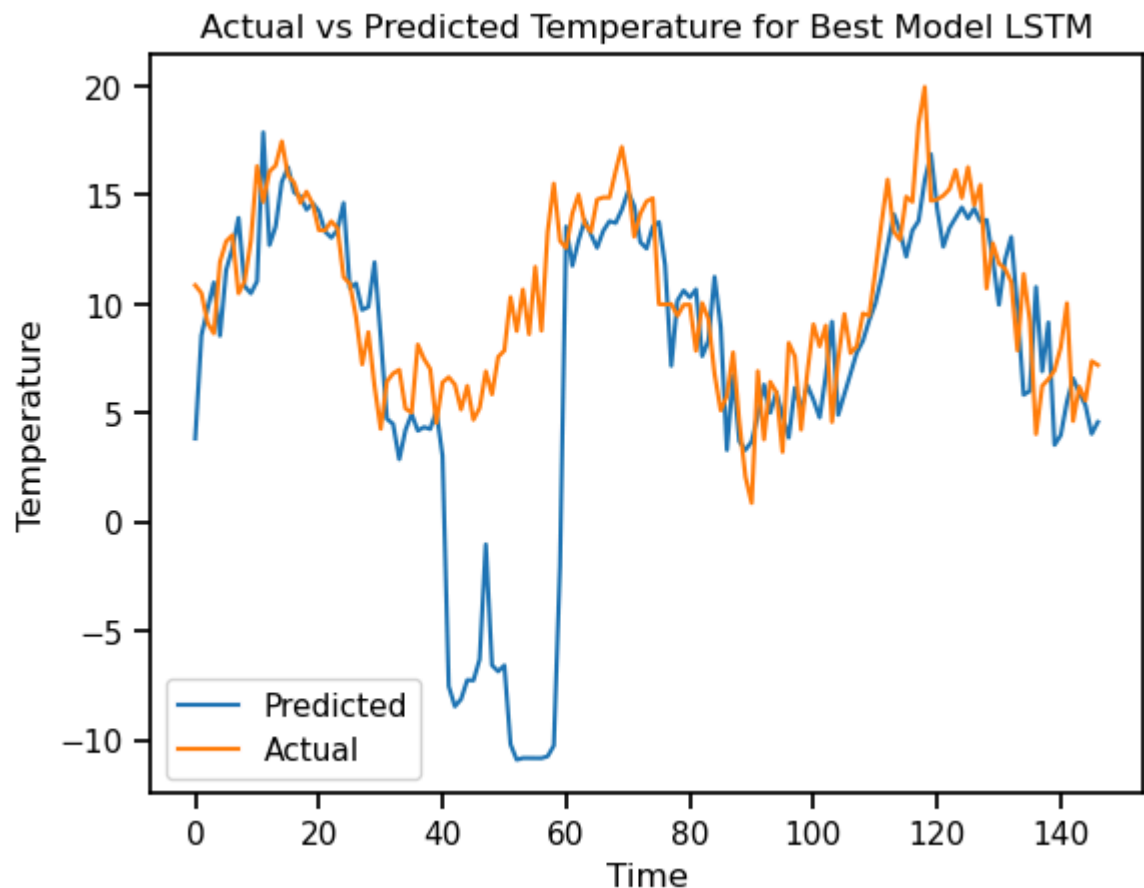
best_model_lstm_temp.hidden = (torch.zeros(1, 1, best_model_lstm_temp.hidden_layer_size),
                                torch.zeros(1, 1, best_model_lstm_temp.hidden_layer_size))
best_test_predictions_lstm_temp.append(best_model_lstm_temp(seq).item())

# Inverse transform the predictions
best_test_predictions_lstm_temp = scaler.inverse_transform(np.array(best_test_predictions_lstm_temp))

# Plot actual vs predicted for the best model
plt.plot(best_test_predictions_lstm_temp, label='Predicted')
plt.plot(scaler.inverse_transform(test_y.reshape(-1, 1)), label='Actual')
plt.xlabel('Time')
plt.ylabel('Temperature')
plt.title('Actual vs Predicted Temperature for Best Model LSTM')
plt.legend()
plt.show()

```

Best model: lr=0.01, hidden\_size=10, epoch: 1 train\_loss\_temperature: 0.03693691  
 val\_loss\_temperature: 0.04347377  
 Best model: lr=0.01, hidden\_size=10, epoch: 26 train\_loss\_temperature: 0.03041685  
 val\_loss\_temperature: 0.02793481  
 Best model: lr=0.01, hidden\_size=10, epoch: 51 train\_loss\_temperature: 0.02838734  
 val\_loss\_temperature: 0.02899118  
 Best model: lr=0.01, hidden\_size=10, epoch: 76 train\_loss\_temperature: 0.02470626  
 val\_loss\_temperature: 0.02907690



## Saving the Model file - LSTM-Temperature

```

In [19]: #Saving the model
# Save the best model as a joblib file
dump(best_model_lstm_temp, 'best_model_lstm_temp.joblib')

# Load the best model from the joblib file ( if you need loading of the model - con
#loaded_model_lstm_temp = load('best_model_lstm_temp.joblib')

```

Out[19]: ['best\_model\_lstm\_temp.joblib']

# SVM

## SVM - Rainfall Forecasting

```
In [20]: import numpy as np
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming you have your data loaded in 'weekly_da'

# Extract rainfall data
rainfall_data = weekly_da['rain'].values.astype(float)

# Normalize data
scaler = MinMaxScaler(feature_range=(-1, 1))
rainfall_data_normalized = scaler.fit_transform(rainfall_data.reshape(-1, 1))

# Define sequence length
seq_length = 5

# Create sequences
def create_sequences(data, seq_length):
    xs = []
    ys = []
    for i in range(len(data) - seq_length):
        x = data[i:i+seq_length]
        y = data[i+seq_length]
        xs.append(x)
        ys.append(y)
    return np.array(xs), np.array(ys)

xs, ys = create_sequences(rainfall_data_normalized, seq_length)

# Split data into training and testing sets
train_size = int(len(xs) * 0.8)
test_size = len(xs) - train_size
train_X, test_X = xs[:train_size], xs[train_size:]
train_y, test_y = ys[:train_size], ys[train_size:]

# Flatten the sequence dimension while keeping the other dimensions intact
train_X_flat = train_X.reshape(train_X.shape[0], -1)
test_X_flat = test_X.reshape(test_X.shape[0], -1)

# Define the parameter grid
param_grid = {
    'kernel': ['rbf', 'linear'], # Different kernel types
    'C': [0.1, 1, 10], # Regularization parameter values
    'gamma': ['scale', 'auto'] # Gamma parameter for RBF kernel
}

# Initialize SVR model
svm_model = SVR()
```

```

# Perform grid search with cross-validation
grid_search = GridSearchCV(estimator=svm_model, param_grid=param_grid, cv=3, scoring='neg_mean_squared_error')
grid_search.fit(train_X_flat, train_y)

# Get the best model
best_svm_model = grid_search.best_estimator_

# Make predictions on the test set using the best model
svm_predictions = best_svm_model.predict(test_X_flat)

# Inverse transform the predictions
svm_predictions_inv = scaler.inverse_transform(svm_predictions.reshape(-1, 1))
test_y_inv = scaler.inverse_transform(test_y.reshape(-1, 1))

# Calculate Mean Squared Error (MSE) for rainfall
mse_rainfall_svm = mean_squared_error(test_y_inv, svm_predictions_inv)
print("MSE for Rainfall (SVM):", mse_rainfall_svm)

# Plot predictions vs actual rainfall
plt.plot(test_y_inv, label='Actual Rainfall')
plt.plot(svm_predictions_inv, label='Predicted Rainfall (SVM)')
plt.xlabel('Time')
plt.ylabel('Rainfall')
plt.title('Actual vs Predicted Rainfall (SVM)')
plt.legend()
plt.show()

# Get the best hyperparameters
best_params = grid_search.best_params_
print("Best Hyperparameters:", best_params)

# Extract grid search results
results = pd.DataFrame(grid_search.cv_results_)

# Pivot the dataframe to have hyperparameters as columns
pivot_table = results.pivot_table(index='param_kernel', columns='param_C', values='mean_test_score')

# Plot the heatmap
plt.figure(figsize=(4, 2))
sns.heatmap(pivot_table, annot=True, cmap="YlGnBu", cbar_kws={'label': 'Mean Test Score'})
plt.title("Grid Search Results")
plt.xlabel("C (Regularization Parameter)")
plt.ylabel("Kernel")
plt.show()

# Create a DataFrame for hyperparameters, validation losses, and training times
svm_data = []
for kernel in param_grid['kernel']:
    for C in param_grid['C']:
        for gamma in param_grid['gamma']:
            idx = param_grid['kernel'].index(kernel) * len(param_grid['C']) * len(param_grid['gamma']) + \
                  param_grid['C'].index(C) * len(param_grid['gamma']) + param_grid['gamma'].index(gamma)
            svm_data.append({
                'Kernel': kernel,
                'C': C,
                'Gamma': gamma,
                'Validation Loss': grid_search.cv_results_['mean_test_score'][idx],
                'Training Time': grid_search.cv_results_['mean_fit_time'][idx]
            })

df_hyperparameters_svm = pd.DataFrame(svm_data)

```

```
# Display the DataFrame  
print(df_hyperparameters_svm)
```

[illegible]



[illegible]

```
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```


```
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```


```
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
MSE for Rainfall (SVM): 3602.8760644338677
```



```
Best Hyperparameters: {'C': 1, 'gamma': 'auto', 'kernel': 'rbf'}
```



	Kernel	C	Gamma	Validation Loss	Training Time
0	rbf	0.1	scale	-0.084488	0.007221
1	rbf	0.1	auto	-0.077342	0.005333
2	rbf	1.0	scale	-0.079407	0.005712
3	rbf	1.0	auto	-0.077342	0.004985
4	rbf	10.0	scale	-0.087609	0.009035
5	rbf	10.0	auto	-0.076622	0.007643
6	linear	0.1	scale	-0.076276	0.007269
7	linear	0.1	auto	-0.076622	0.007356
8	linear	1.0	scale	-0.100665	0.020695
9	linear	1.0	auto	-0.076409	0.014989
10	linear	10.0	scale	-0.078939	0.011009
11	linear	10.0	auto	-0.076409	0.015631

```
In [21]: #This line of code is generated to draw the hyperparameter comparison (using a df to
#df_hyperparameters_svm.to_excel('df_hyperparameters_svm_rain.xlsx'))
```

## SVM - Temperature Forecasting

```
In [22]: # Extract temperature data
temperature_data = weekly_da['temp'].values.astype(float)

# Normalize data
scaler_temp = MinMaxScaler(feature_range=(-1, 1))
temperature_data_normalized = scaler_temp.fit_transform(temperature_data.reshape(-1, 1))

# Create sequences for temperature
xs_temp, ys_temp = create_sequences(temperature_data_normalized, seq_length)

# Split data into training and testing sets for temperature
train_size_temp = int(len(xs_temp) * 0.8)
test_size_temp = len(xs_temp) - train_size_temp
train_X_temp, test_X_temp = xs_temp[:train_size_temp], xs_temp[train_size_temp:]
train_y_temp, test_y_temp = ys_temp[:train_size_temp], ys_temp[train_size_temp:]

# Flatten the sequence dimension while keeping the other dimensions intact for temp
train_X_flat_temp = train_X_temp.reshape(train_X_temp.shape[0], -1)
test_X_flat_temp = test_X_temp.reshape(test_X_temp.shape[0], -1)

# Perform grid search with cross-validation for temperature
grid_search_temp = GridSearchCV(estimator=svm_model, param_grid=param_grid, cv=3, scoring='neg_mean_cross_entropy')
grid_search_temp.fit(train_X_flat_temp, train_y_temp)

# Get the best model for temperature
best_svm_model_temp = grid_search_temp.best_estimator_

# Make predictions on the test set using the best model for temperature
```

```

svm_predictions_temp = best_svm_model_temp.predict(test_X_flat_temp)

# Inverse transform the predictions for temperature
svm_predictions_inv_temp = scaler_temp.inverse_transform(svm_predictions_temp.reshape(-1, 1))
test_y_inv_temp = scaler_temp.inverse_transform(test_y_temp.reshape(-1, 1))

# Calculate Mean Squared Error (MSE) for temperature
mse_temperature_svm = mean_squared_error(test_y_inv_temp, svm_predictions_inv_temp)
print("MSE for Temperature (SVM):", mse_temperature_svm)

# Plot predictions vs actual temperature
plt.plot(test_y_inv_temp, label='Actual Temperature')
plt.plot(svm_predictions_inv_temp, label='Predicted Temperature (SVM)')
plt.xlabel('Time')
plt.ylabel('Temperature')
plt.title('Actual vs Predicted Temperature (SVM)')
plt.legend()
plt.show()

# Get the best hyperparameters for temperature
best_params_temp = grid_search_temp.best_params_
print("Best Hyperparameters for Temperature:", best_params_temp)

# Extract grid search results for temperature
results_temp = pd.DataFrame(grid_search_temp.cv_results_)

# Pivot the dataframe to have hyperparameters as columns for temperature
pivot_table_temp = results_temp.pivot_table(index='param_kernel', columns='param_C')

# Plot the heatmap for temperature
plt.figure(figsize=(4, 2))
sns.heatmap(pivot_table_temp, annot=True, cmap="YlGnBu", cbar_kws={'label': 'Mean 1
plt.title("Grid Search Results for Temperature")
plt.xlabel("C (Regularization Parameter)")
plt.ylabel("Kernel")
plt.show()

# Create a DataFrame for hyperparameters, validation losses, and training times for
svm_data_temp = []
for kernel in param_grid['kernel']:
    for C in param_grid['C']:
        for gamma in param_grid['gamma']:
            idx = param_grid['kernel'].index(kernel) * len(param_grid['C']) * len(p
                param_grid['C'].index(C) * len(param_grid['gamma']) + param_grid[
            svm_data_temp.append({
                'Kernel': kernel,
                'C': C,
                'Gamma': gamma,
                'Validation Loss': grid_search_temp.cv_results_['mean_test_score'][
                'Training Time': grid_search_temp.cv_results_['mean_fit_time'][idx]
            })

df_hyperparameters_svm_temp = pd.DataFrame(svm_data_temp)

# Display the DataFrame for temperature
print(df_hyperparameters_svm_temp)

#This line of code is generated to draw the hyperparameter comparison (using a df t
#df_hyperparameters_svm_temp.to_excel('df_hyperparameters_svm_temp.xlsx')

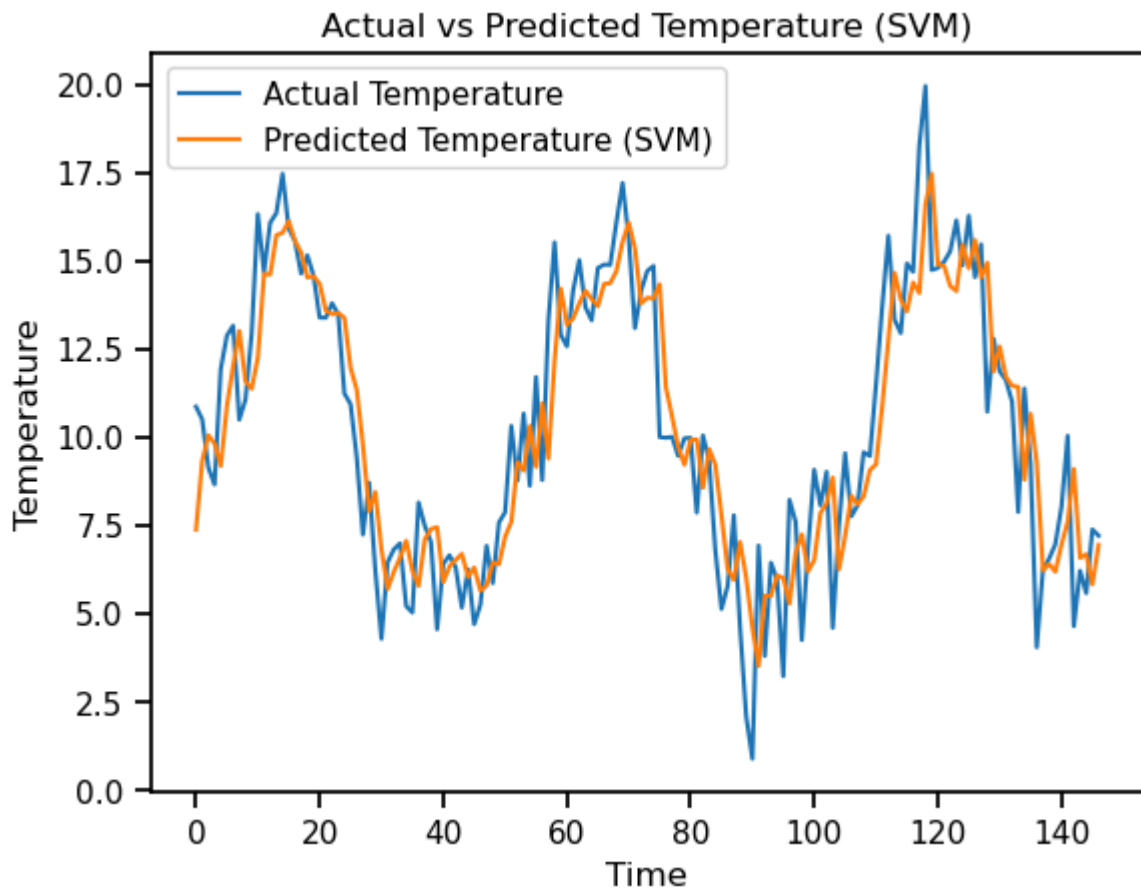
```

[illegible]

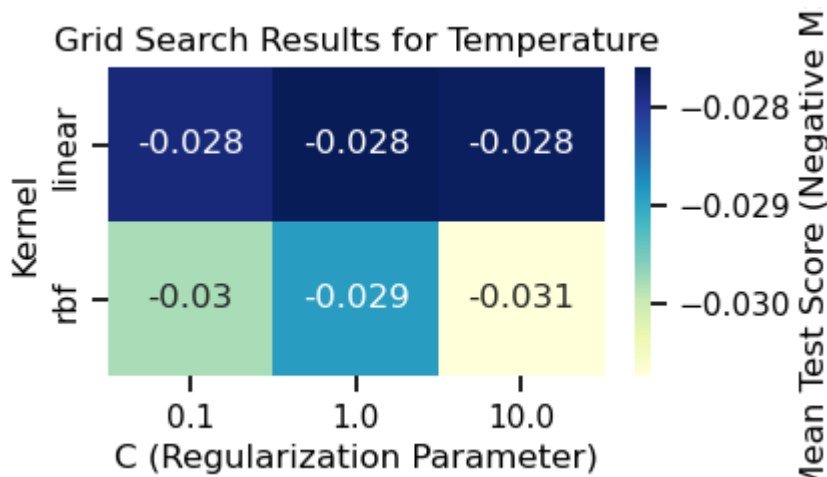


[illegible]

```
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
C:\Anaconda\Lib\site-packages\sklearn\utils\validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
MSE for Temperature (SVM): 3.5063773157897056
```



Best Hyperparameters for Temperature: {'C': 1, 'gamma': 'auto', 'kernel': 'rbf'}



	Kernel	C	Gamma	Validation Loss	Training Time
0	rbf	0.1	scale	-0.031030	0.004668
1	rbf	0.1	auto	-0.027808	0.003674
2	rbf	1.0	scale	-0.028554	0.004379
3	rbf	1.0	auto	-0.027808	0.003957
4	rbf	10.0	scale	-0.030811	0.007997
5	rbf	10.0	auto	-0.027592	0.008041
6	linear	0.1	scale	-0.026985	0.005362
7	linear	0.1	auto	-0.027592	0.007663
8	linear	1.0	scale	-0.033502	0.015662
9	linear	1.0	auto	-0.027631	0.014617
10	linear	10.0	scale	-0.028005	0.009333
11	linear	10.0	auto	-0.027631	0.017263

## Saving the Model files for SVM - Rainfall & Temperature

```
In [23]: # Save the best SVM model for rainfall prediction
dump(best_svm_model, 'best_svm_model_rainfall.joblib')

# Save the best SVM model for temperature prediction
dump(best_svm_model_temp, 'best_svm_model_temperature.joblib')
```

```
Out[23]: ['best_svm_model_temperature.joblib']
```

## Results and Findings - SVM VS LSTM

### MSE PLOTS - Rainfall and Temperature

```
In [24]: # Calculate Mean Absolute Error (MAE) for rainfall predictions
mae_rainfall_svm = mean_absolute_error(test_y_inv, svm_predictions_inv)
mae_rainfall_lstm = mean_absolute_error(test_y_inv, best_test_predictions_rain_lstm)

# Calculate Mean Absolute Error (MAE) for temperature predictions
mae_temperature_svm = mean_absolute_error(test_y_inv_temp, svm_predictions_inv_temp)
mae_temperature_lstm = mean_absolute_error(test_y_inv_temp, best_test_predictions_temp_lstm)

# Calculate the MSE values for SVM models for temperature and rainfall predictions
lowest_temp_svm = mse_temperature_svm
lowest_rain_svm = mse_rainfall_svm

# Calculate the MSE values for LSTM models for temperature and rainfall predictions
lowest_temp_lstm = mse_df_temp_lstm['mse_temperature'].min()
lowest_rain_lstm = mse_df_rain_lstm['mse_rainfall'].min()

# Plotting subplots for MSE and MAE comparison for temperature and rainfall
```



```

# Create a figure and subplots
fig, axs = plt.subplots(2, 2, figsize=(12, 10))

# Temperature MSE comparison
axs[0, 0].bar(['SVM', 'LSTM'], [lowest_temp_svm, lowest_temp_lstm], color=['orange', 'green'])
axs[0, 0].set_xlabel('Model')
axs[0, 0].set_ylabel('MSE')
axs[0, 0].set_title('Temperature MSE (SVM vs LSTM)')
# Annotate values
for i, v in enumerate([lowest_temp_svm, lowest_temp_lstm]):
    axs[0, 0].text(i, v + 0.01, str(round(v, 2)), ha='center', va='bottom')

# Rainfall MSE comparison
axs[0, 1].bar(['SVM', 'LSTM'], [lowest_rain_svm, lowest_rain_lstm], color=['orange', 'green'])
axs[0, 1].set_xlabel('Model')
axs[0, 1].set_ylabel('MSE')
axs[0, 1].set_title('Rainfall MSE (SVM vs LSTM)')
# Annotate values
for i, v in enumerate([lowest_rain_svm, lowest_rain_lstm]):
    axs[0, 1].text(i, v + 0.01, str(round(v, 2)), ha='center', va='bottom')


# Temperature MAE comparison
axs[1, 0].bar(['SVM', 'LSTM'], [mae_temperature_svm, mae_temperature_lstm], color=['orange', 'green'])
axs[1, 0].set_xlabel('Model')
axs[1, 0].set_ylabel('MAE')
axs[1, 0].set_title('Temperature MAE (SVM vs LSTM)')
# Annotate values
for i, v in enumerate([mae_temperature_svm, mae_temperature_lstm]):
    axs[1, 0].text(i, v + 0.01, str(round(v, 2)), ha='center', va='bottom')

# Rainfall MAE comparison
axs[1, 1].bar(['SVM', 'LSTM'], [mae_rainfall_svm, mae_rainfall_lstm], color=['orange', 'green'])
axs[1, 1].set_xlabel('Model')
axs[1, 1].set_ylabel('MAE')
axs[1, 1].set_title('Rainfall MAE (SVM vs LSTM)')
# Annotate values
for i, v in enumerate([mae_rainfall_svm, mae_rainfall_lstm]):
    axs[1, 1].text(i, v + 0.01, str(round(v, 2)), ha='center', va='bottom')

# Adjust layout
plt.tight_layout()

# Show plot
plt.show()

```



## Residual plots - SVM VS LSTM

```
In [25]: # Calculate residuals for SVM temperature prediction
residuals_svm_temp = svm_predictions_temp - test_y_inv_temp.flatten()
residuals_svm_rain = svm_predictions_inv.flatten() - test_y_inv.flatten()
# Calculate residuals for LSTM temperature prediction
best_test_predictions_lstm_temp = best_test_predictions_lstm_temp.flatten()
residuals_lstm_temp = best_test_predictions_lstm_temp - test_y_inv_temp.flatten()
best_test_predictions_rain_lstm = best_test_predictions_rain_lstm.flatten()
test_y_inv_rain = scaler.inverse_transform(test_y.reshape(-1, 1))
residuals_rain_lstm = best_test_predictions_rain_lstm.flatten() - test_y_inv_rain.flatten()
# Create a figure with subplots for rainfall and temperature predictions
plt.figure(figsize=(10, 8))
# Subplot for rainfall predictions
plt.subplot(2, 1, 1)

# Residual plot for SVM rainfall prediction
plt.scatter(svm_predictions_inv, residuals_svm_rain, color='orange', alpha=0.5, label='SVM')
plt.axhline(y=0, color='red', linestyle='--')
plt.title('Residual Plot for Rainfall Prediction')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.grid(True)
plt.legend()

# Residual plot for LSTM rainfall prediction
plt.scatter(best_test_predictions_rain_lstm, residuals_rain_lstm, color='blue', alpha=0.5, label='LSTM')
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.grid(True)
```

```

plt.legend()


# Subplot for temperature predictions
plt.subplot(2, 1, 2)

# Residual plot for SVM temperature prediction
plt.scatter(svm_predictions_temp, residuals_svm_temp, color='orange', alpha=0.5, label='SVM')
plt.axhline(y=0, color='red', linestyle='--')
plt.title('Residual Plot for Temperature Prediction')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.grid(True)
plt.legend()

# Residual plot for LSTM temperature prediction
plt.scatter(best_test_predictions_lstm_temp, residuals_lstm_temp, color='blue', alpha=0.5, label='LSTM')
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()

```



In [ ]: